

Quality-cost trade-off across composition methods (Qwen-2.5-0.5B)

Quality (aggregate KL, aligned tokenization)

Cost (Table 3)

full prefill
345 ms

M1/M2 baseline

